

Process Management

Deadlock

Introduction

- In a multiprogramming environment, several processes may compete for a finite number of resources.
- A process requests resources; if the resources are not available at that time, the process enters a waiting state. Sometimes, a waiting process is never again able to change state, because the resources it has requested are held by other waiting processes, blocking the process.

Introduction


- Deadlock is a situation in which a set of blocked processes, each holding a resource, is waiting to acquire a resource held by another process in the set.
- Example
 - System has 2 tape drives.
 - P_1 and P_2 each hold one tape drive and each needs another one.
- Example
 - Semaphores A and B , initialized to 1

P_0	P_1
<i>wait (A);</i>	<i>wait(B)</i>
<i>wait (B);</i>	<i>wait(A)</i>

Dr. Padma D. Adane
Department of IT, RCOEM

3

Traffic Gridlock



- Deadlock can be resolved if one car backs up (preempt resources and rollback).
- Several cars may have to be backed up if a deadlock occurs; Starvation is possible.

Dr. Padma D. Adane
Department of IT, RCOEM

4

System Model

- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- Each resource type R_i has W_i instances.
- Each process utilizes a resource as follows:
 - request
 - use
 - release

Dr. Padma D. Adane
Department of IT, RCOEM

5

Deadlock Characterization

- Deadlock can arise if four conditions hold simultaneously.
 - **Mutual exclusion:** only one process at a time can use a resource.
 - **Hold and wait:** a process holding at least one resource is waiting to acquire additional resources held by other processes.
 - **No preemption:** a resource can be released only voluntarily by the process holding it, after that process has completed its task.
 - **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_n is waiting for a resource that is held by P_0 .

Dr. Padma D. Adane
Department of IT, RCOEM

6



Resource-Allocation Graph



- A set of vertices V and a set of edges E .
- V is partitioned into two types:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system.
- Request edge – directed edge $P_i \rightarrow R_j$
- Assignment edge – directed edge $R_j \rightarrow P_i$

Dr. Padma D. Adane
Department of IT, RCOEM

7

Resource-Allocation Graph

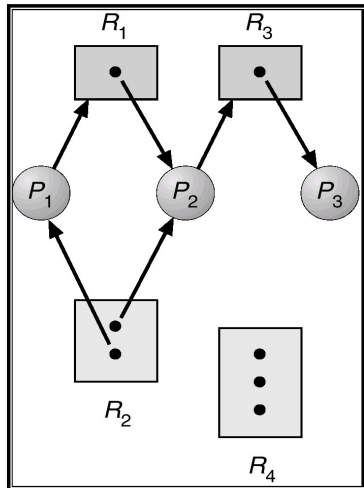
- Process 
- Resource Type with 4 instances 
- P_i requests instance of R_j
- P_i is holding an instance of R_j



Dr. Padma D. Adane
Department of IT, RCOEM

8

Example of a Resource Allocation Graph



- Process P_1 is holding an instance of resource R_2 and is waiting for an instance of resource type R_1 .
- Process P_2 is holding an instance of R_1 and an instance of R_2 and is waiting for an instance of R_3 .
- Process P_3 is holding an instance of R_3

Dr. Padma D. Adane
Department of IT, RCOEM

9


Resource-Allocation Graph

- If a graph contains no cycles, then no process is in deadlock. However, if the graph contains a cycle, then a deadlock **MAY** exist.
- If each resource type has exactly one instance, then a cycle implies that a deadlock has occurred.
- If the cycle involves only a set of resource types, each of which has only a single instance, then a deadlock has occurred. Each process involved in the cycle is deadlocked.
- In this case, a cycle in the graph is a necessary and sufficient condition for the existence of deadlock.

Dr. Padma D. Adane
Department of IT, RCOEM

10

If each resource type has several instances, then a cycle does not necessary imply that a deadlock has occurred. In this case, a cycle in the graph is a necessary but not a sufficient condition for the existence of deadlock




$P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$

$P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$

Dr. Padma D. Adane
Department of IT, RCOEM

11



$P_1 \rightarrow R_1 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$

The system is not in deadlock. P_4 may release its instance of resource type R_2 which can be allocated to P_3 , breaking the cycle.

Dr. Padma D. Adane
Department of IT, RCOEM

12

Methods for Handling Deadlocks

- Ensure that the system will **never** enter a deadlock state.
 - Deadlock Prevention
 - Deadlock Avoidance
- Allow the system to enter a deadlock state and then recover.
 - Deadlock Detection and Recovery
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX/Linux.

Dr. Padma D. Adane
Department of IT, RCOEM

13

Deadlock Prevention

- Ensure that one of the necessary conditions for deadlocks does not hold. Can be accomplished by restraining the ways requests can be made.
- **Mutual Exclusion:** Must hold for non-sharable resources that can be accessed simultaneously by various processes. For example, a printer cannot be simultaneously shared by several processes.
- Certain resources are intrinsically non sharable. Therefore, Mutual exclusion cannot be used for deadlock prevention.

Dr. Padma D. Adane
Department of IT, RCOEM

14

Deadlock Prevention

- **Hold and Wait:** Must guarantee that whenever a process request a resource, it does not hold any other resources.
 - One way is to require process to request and be allocated all its resources before its begin its execution. For example, to copy a file from drive to disk, sort it and print it, drive, disk and printer all three needs to be acquired, although printer is required only at the end of the process.
 - Another way is to acquire the resource, use it and then release it before requesting any other resource. Thus, acquire drive and disk, copy and release. Again request disk and printer, print the file and then release both the resources.
 - Both protocols have low resource utilization; starvation problem

Dr. Padma D. Adane
Department of IT, RCOEM

15

Deadlock Prevention

- **No Preemption: Allow Preemption**
 - If a process, that is holding some resources, requests another resource that cannot be immediately allocated to it, then resources currently being held are released.
 - Preempted resources are added to the list of resources for which the process is waiting.
 - Process will be restarted only when it can regain its old resources as well as the new ones that it is requesting
 - Can be applied only to resources (CPU, memory space) whose state can be easily saved restored later.

Dr. Padma D. Adane
Department of IT, RCOEM

16

Deadlock Prevention

- Circular Wait: Impose a total ordering of all resource types.
- A one to one function can be defined as $F: R \rightarrow N$, where N is the set of natural numbers. $F(R1) = n1, F(R2) = n2, \dots$
- Each process can requests resources in an increasing order of enumeration i.e. each process can initially request any number of instances of a resource type, say R_i . After that, the process can request instances of R_j , if and only if, $F(R_j) > F(R_i)$

Dr. Padma D. Adane
Department of IT, RCOEM

17

Deadlock Avoidance

- Requires that the system has some additional *a priori* information available.
- Simplest and most useful model requires that each process declare the *maximum number* of resources of each type that it may need.
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
- Resource-allocation state is defined by the number of available and allocated resources, and the maximum demands of the processes.

Dr. Padma D. Adane
Department of IT, RCOEM

18


Safe State of the System

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a **safe state**.
- System is in **safe state** if there exists a safe sequence of all processes. $\langle P_1, P_2, \dots, P_n \rangle$ such that for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$.
 - If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished.
 - When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate.
 - When P_i terminates, P_{i+1} can obtain its needed resources, and so on.

Dr. Padma D. Adane
Department of IT, RCOEM

19

Deadlock Avoidance



If a system is in safe state \Rightarrow no deadlocks.

If a system is in unsafe state \Rightarrow possibility of deadlock.

Avoidance \Rightarrow ensure that a system will never enter an unsafe state.

Dr. Padma D. Adane
Department of IT, RCOEM

20

Deadlock Avoidance Algorithms

- Single Instance of each resource type
 - Use variant of resource allocation graph
- Multiple instances of each resource type
 - Use Banker's algorithm

Resource Allocation graph with claim edge

- *Claim edge* $P_i \rightarrow R_j$ indicated that process P_i may request resource R_j ; represented by a dashed line.
- Claim edge converts to request edge when a process requests a resource.
- Request edge converts to an assignment edge when the resource is allocated to the process.
- When a resource is released by a process, assignment edge reconverts to a claim edge.
- Resources must be claimed *a priori* in the system.
- A cycle in the graph implies that the system is in unsafe state.


Resource Allocation Graph Algorithm

- Suppose Process P_i requests a resource R_j
- The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the RAG.
- Otherwise, the process must wait till system is safe.

Dr. Padma D. Adane
Department of IT, RCOEM

23

Example




- P1 is holding resource R1 and has a claim on R2
- P2 is requesting R1 and has claim on R2
- Currently, there is no cycle and hence system is in a safe state.

Dr. Padma D. Adane
Department of IT, RCOEM

24

Example



Suppose the request from P2 to R2 is granted and converted to an assignment edge. This will now result in a cycle taking the system to an unsafe state.

Dr. Padma D. Adane
Department of IT, RCOEM

25

Banker's Algorithm

- Let n = number of processes, and m = number of resources types

Data Structures

- Available:** Vector of length m . If available $[j] = k$, there are k instances of resource type R_j available.
- Max:** $n \times m$ matrix. If $Max[i, j] = k$, then process P_i may request at most k instances of resource type R_j .
- Allocation:** $n \times m$ matrix. If $Allocation[i, j] = k$ then P_i is currently allocated k instances of R_j .
- Need:** $n \times m$ matrix. If $Need[i, j] = k$, then P_i may need k more instances of R_j to complete its task.

$$Need[i, j] = Max[i, j] - Allocation[i, j].$$

Dr. Padma D. Adane
Department of IT, RCOEM

26

Banker's Algorithm

Request = request vector for process P_i . If $Request_i[j] = k$ then process P_i wants k instances of resource type R_j .

1. If $Request_i \leq Need_i$ go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
2. If $Request_i \leq Available$, go to step 3. Otherwise P_i must wait, since resources are not available.
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

Call Safety algorithm.

- If safe \Rightarrow the resources are allocated to P_i .
- If unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Dr. Padma D. Adane
Department of IT, RCOEM

27

Safety Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively. Initialize:
 $Work = Available$
 $Finish[i] = false$ for $i = 1, 2, \dots, n$.
2. Find an i such that both:
 (a) $Finish[i] = false$
 (b) $Need_i \leq Work$
 If no such i exists, go to step 4.
3. $Work = Work + Allocation_i$
 $Finish[i] = true$
 go to step 2.
4. If $Finish[i] == true$ for all i , then the system is in a safe state.

Dr. Padma D. Adane
Department of IT, RCOEM

28

Example

- Let there be 5 processes P_0 through P_4 in the system and 3 resource types A with 10 instances, B with 5 instances and C with 7 instances
- At time t_0 the system state is as follows:

	<i>Allocation</i>			<i>Max</i>			<i>Available</i>		
	<i>A</i>	<i>B</i>	<i>C</i>	<i>A</i>	<i>B</i>	<i>C</i>	<i>A</i>	<i>B</i>	<i>C</i>
P_0	0	1	0	7	5	3	3	3	2
P_1	2	0	0	3	2	2			
P_2	3	0	2	9	0	2			
P_3	2	1	1	2	2	2			
P_4	0	0	2	4	3	3			

Dr. Padma D. Adane
Department of IT, RCOEM

29

- From the given data, calculate the Need matrix as $\text{Need}[i] = \text{Max}[i] - \text{Allocation}[i]$

	<i>Need</i>		
	<i>A</i>	<i>B</i>	<i>C</i>
P_0	7	4	3
P_1	1	2	2
P_2	6	0	0
P_3	0	1	1
P_4	4	3	1

Is the system in a safe state?

Dr. Padma D. Adane
Department of IT, RCOEM

30

$$\text{Work} = [3 \ 3 \ 2]$$

$$\text{Finish} =$$

0	1	2	3	4
F	F	F	F	F

$$\text{Need}_0 [7 \ 4 \ 3] \geq \text{Work}$$

$$\text{Need}_1 [1 \ 2 \ 2] \leq \text{Work}; \text{Work} = [3 \ 3 \ 2] + [2 \ 0 \ 0] \\ = [5 \ 3 \ 2]$$

$$\text{Finish} =$$

0	1	2	3	4
F	T	F	F	F

$$\text{Need}_2 [6 \ 0 \ 0] \geq \text{Work};$$

$$\text{Need}_3 [0 \ 1 \ 1] \leq \text{Work}; \text{Work} = [5 \ 3 \ 2] + [2 \ 1 \ 1] \\ = [7 \ 4 \ 3]$$

Dr. Padma D. Adane
Department of IT, RCOEM

31

$$\text{Finish} =$$

0	1	2	3	4
F	T	F	T	F

$$\text{Need}_4 [4 \ 3 \ 1] \leq \text{Work}; \text{Work} = [7 \ 4 \ 3] + [0 \ 0 \ 2] \\ = [7 \ 4 \ 5]$$

$$\text{Finish} =$$

0	1	2	3	4
F	T	F	T	T

$$\text{Need}_0 [7 \ 4 \ 3] \leq \text{Work}; \text{Work} = [7 \ 4 \ 5] + [0 \ 1 \ 0] \\ = [7 \ 5 \ 5]$$

$$\text{Finish} =$$

0	1	2	3	4
T	T	F	T	T

Dr. Padma D. Adane
Department of IT, RCOEM

32

$$\text{Need}_2 [6\ 0\ 0] \leq \text{Work}; \text{Work} = [7\ 5\ 5] + [3\ 0\ 2] \\ = [10\ 5\ 7]$$

Finish =	0	1	2	3	4
	T	T	T	T	T

The system is safe

Safe Sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$

There can be more than one safe sequence

Example

- Consider again the safe state of the system

	<i>Allocation</i>	<i>Max</i>	<i>Available</i>
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	

- Can the request from Process P_1 for $[1\ 0\ 2]$ be granted?

	Need		
	A	B	C
P_0	7	4	3
P_1	1	2	2
P_2	6	0	0
P_3	0	1	1
P_4	4	3	1

Dr. Padma D. Adane
Department of IT, RCOEM

35

- Is $Request_1 \leq Need_1$ $[1\ 2\ 2] \leq [1\ 0\ 2]$
- Is $request_1 \leq Available$ $[1\ 2\ 2] \leq [3\ 3\ 2]$
- $Available = Available - Request_1$
 - $Available = [3\ 3\ 2] - [1\ 0\ 2] = [2\ 3\ 0]$
- $Allocation_1 = Allocation_1 + Request_1$
 - $Allocation_1 = [2\ 0\ 0] + [1\ 0\ 2] = [3\ 0\ 2]$
- $Need_1 = Need_1 - Request_1$
 - $Need_1 = [1\ 2\ 2] - [1\ 0\ 2] = [0\ 2\ 0]$

Dr. Padma D. Adane
Department of IT, RCOEM

36

	<i>Allocation</i>	<i>Need</i>	<i>Available</i>
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 1	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

Check if the system is safe.

Sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement.

Dr. Padma D. Adane
Department of IT, RCOEM

37

Additional Requests

- Can request for (3,3,0) by P_4 be granted?
- Can request for (0,2,0) by P_0 be granted?

Dr. Padma D. Adane
Department of IT, RCOEM

38

Deadlock Detection

- Certain systems do not take precautionary steps and let deadlock happen; it runs deadlock detection algorithm periodically to detect it and then recover from it.
- For all resources having only a single instance, detection algorithm uses a variant of resource allocation graph, called as wait-for graph.
- For resources having several instances, the algorithm uses several time varying data structures similar to the banker's algorithm.

Dr. Padma D. Adane
Department of IT, RCOEM

39

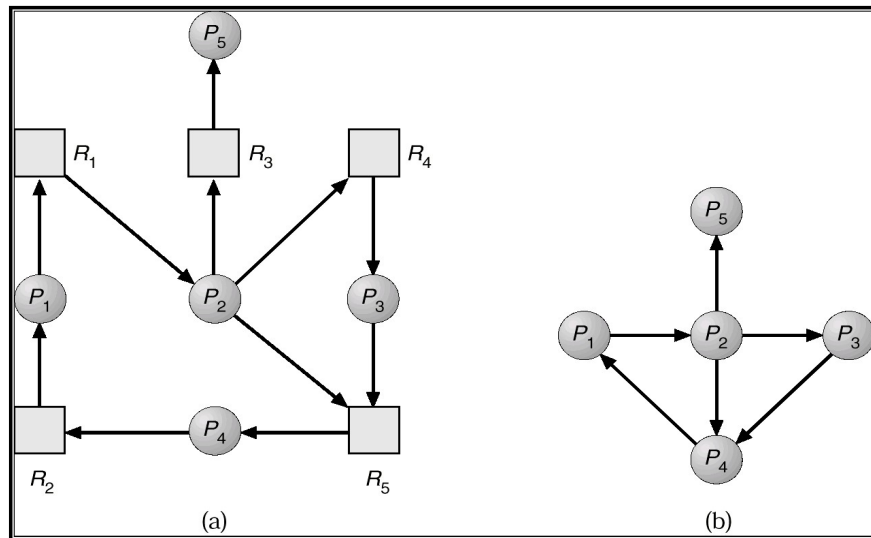
Single Instance of Each Resource Type

- In a wait-for graph, an edge from P_i to P_j implies that process P_i is waiting for a resource held by process P_j .
- An edge $P_i \rightarrow P_j$ exists in a wait-for graph if and only if the corresponding Resource-allocation graph has two edges $P_i \rightarrow R_q$ and $R_q \rightarrow P_j$ for some resource R_q .
- A deadlock exists in the system if and only if the wait-for graph contains a cycle.
- An algorithm to detect a cycle in a graph requires an order of n^2 operations, where n is the number of vertices in the graph.

Dr. Padma D. Adane
Department of IT, RCOEM

40

Resource allocation graph and equivalent Wait-for graph



Dr. Padma D. Adane
Department of IT, RCOEM

41

Several Instances of a Resource Type

■ Data structures used

- **Available:** A vector of length m indicating the number of available resources of each type.
- **Allocation:** An $n \times m$ matrix defining the number of resources of each type currently allocated to each process.
- **Request:** An $n \times m$ matrix indicating the current request of each process. If $Request[i_j] = k$, then process P_i is requesting k more instances of resource type R_j .

Dr. Padma D. Adane
Department of IT, RCOEM

42

Deadlock Detection Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively Initialize:

(a) *Work* = *Available*

(b) For $i = 1, 2, \dots, n$, if $Allocation_i \neq 0$, then $Finish[i] = false$; otherwise, $Finish[i] = true$.

2. Find an index *i* such that both:

(a) $Finish[i] == false$

(b) $Request_i \leq Work$

If no such *i* exists, go to step 4.

Dr. Padma D. Adane
Department of IT, RCOEM

43

Deadlock Detection Algorithm

3 $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2.

4 If $Finish[i] == false$, for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if $Finish[i] == false$, then P_i is deadlocked.

- Algorithm requires an order of $O(m \times n^2)$ operations to detect whether the system is in deadlocked state.

Dr. Padma D. Adane
Department of IT, RCOEM

44

Example

- Five processes P_0 through P_4 ; three resource types A (7 instances), B (2 instances), and C (6 instances).
- Snapshot at time T_0 :

	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	0	0	0	0	0	0
P_1	2	0	0	2	0	2			
P_2	3	0	3	0	0	0			
P_3	2	1	1	1	0	0			
P_4	0	0	2	0	0	2			

- Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in $Finish[i] = \text{true}$ for all i .

Dr. Padma D. Adane
Department of IT, RCOEM

45

Example

- Let P_2 requests an additional instance of type C

	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	0	0	0	0	0	0
P_1	2	0	0	2	0	2			
P_2	3	0	3	0	0	1			
P_3	2	1	1	1	0	0			
P_4	0	0	2	0	0	2			

Is the system in deadlock?

- Deadlock exists and processes P_1 , P_2 , P_3 , and P_4 are in deadlock

Dr. Padma D. Adane
Department of IT, RCOEM

46

Detection Algorithm Usage

- When, and how often, to invoke depends on:
 - How often a deadlock is likely to occur?
 - How many processes will need to be rolled back?
 - one for each disjoint cycle
- If detection algorithm is invoked after every resource request that cannot be granted immediately, it is possible to find the process leading to deadlock along with the processes involved in the deadlock; however, this results in considerable overhead in computation time.
- If detection algorithm is invoked arbitrarily, say, once every hour or whenever CPU utilization falls below 40%, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock.

Dr. Padma D. Adane
Department of IT, RCOEM

47

System Recovery from Deadlock

- Two Alternatives
 - Process Termination
 - Resource Preemption
- Process Termination
 - Abort all deadlocked processes: An expensive option as the processes may have computed for a long time.
 - Abort one process at a time until the deadlock cycle is eliminated: Incurs lot of overhead as after aborting each process detection algorithm needs to be run to check if the deadlock is over.

Dr. Padma D. Adane
Department of IT, RCOEM

48

System Recovery from Deadlock

- In which order should we choose to abort?
 - ❑ Priority of the process.
 - ❑ How long process has computed, and how much longer to completion.
 - ❑ Resources the process has used.
 - ❑ Resources process needs to complete.
 - ❑ How many processes will need to be terminated.
 - ❑ Is process interactive or batch?

Dr. Padma D. Adane
Department of IT, RCOEM

49

System Recovery from Deadlock

- Resource Preemption: Three issues needs to be addressed while dealing with deadlock using resource preemption
 - ❑ Selecting a victim: Which Process and which resource to preempt? One with minimum cost; in terms of how many resources the process is holding? how much work has been done? how much work is left?
 - ❑ Rollback: Restart process; a costly affair, return to some safe state; requires system to record lot of information about each process.
 - ❑ Starvation: Same process may always be picked as victim; include number of rollback in cost factor.

Dr. Padma D. Adane
Department of IT, RCOEM

50

Combined Approach to Deadlock Handling

- Combine the three basic approaches

- prevention
- avoidance
- detection

allowing the use of the optimal approach for each of resources in the system.


- Partition resources into hierarchically ordered classes.
- Use most appropriate technique for handling deadlocks within each class.

Dr. Padma D. Adane
Department of IT, RCOEM

51

Example

- Is the system, represented by the RAG given below, in deadlock?



- P1 and P3 are deadlocked because a cycle exists between them. P2 is also deadlocked as all resources required by it and held by P3 which will not release them since it is deadlocked.

Dr. Padma D. Adane
Department of IT, RCOEM

52

Example

Consider the following snapshot of a system

	Allocation				Max				Available			
	A	B	C	D	A	B	C	D	A	B	C	D
P0	0	0	1	2	0	0	1	2	1	5	2	0
P1	1	0	0	0	1	7	5	0				
P2	1	3	5	4	2	3	5	6				
P3	0	6	3	2	0	6	5	2				
P4	0	0	1	4	0	6	5	6				

Is the system safe?

Can a request from Request₃[0 4 2 0] be granted?